

EvoStrategy: Document Reconciliation Pipeline and Verified Revenue Intelligence for IT Services Enterprise

Mazin Moosa, Ayushi Rastogi, Prateek Mallappa Hulamani, Adham Shakeer

Department of Artificial Intelligence and Machine Learning

Mangalore Institute of Technology & Engineering, Mangalore, India

Abstract—IT services enterprises generate large volumes of operational and financial information across heterogeneous documents, including invoices, project reports, contracts and financial statements. The increasing fragmentation of such information across formats creates challenges in maintaining reliable, consistent, and traceable data for enterprise decision-making. Numerical inconsistencies, entity mismatches, missing records, and contradictory information across documents can compromise the reliability of downstream financial analysis and forecasting. EvoStrategy is a unified enterprise intelligence framework designed to transform fragmented, multi-source business documents into a reconciled and verifiable data foundation that enables reliable revenue forecasting and advanced analytics. It addresses this problem by establishing data reconciliation and verification as a prerequisite for generating trusted enterprise intelligence. The proposed system integrates OCR-based document ingestion, cross-document entity and numerical reconciliation and human-in-the-loop verification through a unified pipeline. Verified records are subsequently used for profit forecasting, analytics and What-If simulation, with complete auditability and local deployment to preserve data integrity and enterprise confidentiality.

I. INTRODUCTION

Modern IT services enterprises generate and maintain large volumes of operational and financial information as part of their daily activities. This information is distributed across different types of business documents, including invoices, purchase orders, project reports, client contracts, payment records, ledgers, and financial statements. These documents are often created at different stages of business operations and may be stored in different formats such as PDF files, spreadsheets, scanned documents and images. As a result, an enterprise may possess a large amount of information about the same transaction or business activity without having a single, consistent representation of that information. The implementation plan for EvoStrategy identifies this fragmentation as a key challenge, particularly because such documents can contain numerical inconsistencies and conflicting information.

The problem becomes more significant when organizations attempt to use these documents for financial analysis and decision-making. Business decisions related to expenditure, budgeting, revenue, profitability, and future financial planning depend heavily on the accuracy of the underlying data. For example, an invoice may contain an amount that differs from the corresponding purchase order or ledger entry. Similarly, the same vendor or client may be represented using different names across documents, making it difficult to determine whether multiple records refer to the same entity. Documents may also be missing, duplicated, or contain incomplete information. These inconsistencies can remain unnoticed when information is processed independently rather than being compared across its related sources.

Traditional document-processing approaches generally focus on extracting information from individual documents. Optical Character Recognition (OCR), for example, can convert information from scanned documents and images into machine-readable text. Although this is an important step toward automation, extracting information alone does not establish that the extracted information is correct or consistent with other available records. A system may successfully extract an amount from an invoice while still failing to determine whether that amount agrees with the corresponding purchase order, payment record, or ledger. Therefore, document extraction and business-data validation represent two different challenges. A reliable enterprise intelligence system must address both.

Another important challenge is the identification of relationships between documents. Financial and operational records are not isolated entities; they are connected through common attributes such as vendors, dates, transaction amounts, document types, and business activities. Establishing these relationships is necessary before meaningful comparisons can be performed. However, exact matching is often insufficient because the same entity can appear under different names or formats. EvoStrategy therefore incorporates entity resolution using techniques such as fuzzy matching and named-entity recognition, together with document-linking rules based on attributes such as vendor similarity and date proximity.

Even after related documents have been identified, the system must determine whether differences between corresponding records represent genuine errors or acceptable variations. A simple comparison of numerical values may incorrectly classify small differences as contradictions. To address this issue, EvoStrategy introduces tolerance-based numerical comparison, where discrepancies within an acceptable range can be automatically resolved while larger discrepancies are escalated for further examination. The implementation plan specifies a default tolerance band of 2% and provides separate escalation logic for discrepancies beyond the tolerance or for missing records. Discrepancies are also classified into categories such as numerical mismatch, missing document, entity alias, and duplicate entry.

A further limitation of fully automated reconciliation is that not every discrepancy can be resolved reliably by an algorithm. Some inconsistencies may require contextual understanding or access to information that is not available to the system. Completely automatic correction could therefore introduce additional errors into financial records. EvoStrategy addresses this limitation through a human-in-the-loop verification stage. The system provides a dashboard where reviewers can inspect source documents and reconciliation results and perform ACCEPT, REJECT, or CORRECT actions. Reviewer identity, timestamps, and correction information are recorded as part of the audit trail. This allows automation to handle routine cases while retaining human judgement for ambiguous or high-risk cases.

The central principle of EvoStrategy is that data reconciliation should occur before analytics and forecasting. This creates a clear separation between raw document

information, reconciled information, human-verified information, and analytical outputs. According to the proposed architecture, the system consists of four stages: document ingestion and OCR, cross-document reconciliation, human-in-the-loop verification and verified analytics and forecasting. This architecture ensures that projections are traceable to the source documents from which the underlying information was obtained.

Therefore, EvoStrategy is designed not simply as a document-processing system, but as an integrated framework for transforming fragmented enterprise documents into a reconciled, verified, and traceable data foundation for financial information. Its primary contribution is the introduction of reconciliation as a necessary stage between document ingestion and analytics. By combining OCR-based extraction, cross-document reconciliation, human verification, audit logging and verified forecasting within a single pipeline, the system seeks to reduce manual reconciliation effort while improving the reliability of downstream financial analysis. The planned evaluation measures this objective through four key outcomes: document extraction accuracy, reconciliation precision and recall, forecasting accuracy, and reduction in human review effort.

II. LITERATURE SURVEY

Table I summarizes representative works surveyed for this project, covering retrieval-augmented generation and knowledge-graph reasoning, document tamper detection, multimodal document understanding, and related machine-learning pipelines. For each work, the table contrasts its methodology and limitations with the research gap it identifies for EvoStrategy.

TABLE I
LITERATURE SURVEY OF EXISTING APPROACHES

S.no	Paper Title	Methodology	Limitations	Research Gaps Identified	Year
1	GraphRAG-R1: Graph Retrieval Augmented Generation with Process Constrained RL	Adaptive GraphRAG with reinforcement learning to improve multi-hop retrieval and reasoning across structured knowledge graphs	Focused on benchmark QA tasks; not evaluated on enterprise document reconciliation or financial simulation use cases	GraphRAG retrieval patterns can be adapted to trace decision outcome chains in IT financial documents for evidence-based budget projections	2025
2	Development of an Offline RAG Chatbot for Answering Food Hygiene and Safety Questions Based on Vietnamese Legal Frameworks	Offline RAG pipeline using Qwen2.5 and Chroma vector database for domain-specific corpus of 170 structured legal documents	Lacks multi-hop reasoning for conflicting document versions; limited to homogeneous corpus of structured statutory text	Offline RAG architecture demonstrates sovereign deployment model directly applicable to enterprise financial document corpora requiring data confidentiality	2026
3	Bridging Legal Knowledge and AI: Retrieval-Augmented Generation with Vector Stores, Knowledge Graphs, and Hierarchical Non-negative Matrix Factorization	Modular agentic RAG with Neo4j knowledge graph for semantic retrieval and structural reasoning across legal statutes and case law	Incomplete entity attribution; limited integration of cross-document reconciliation for numerical contradictions	Hybrid vector + graph retrieval architecture provides blueprint for linking invoice records to payment memos and purchase orders in EvoStrategy	2025
4	A Review on Retrieval-Augmented Generation: Architectures, Research Challenges, and Emerging Frontiers	Analytical synthesis of RAG evolution evaluating structural coupling between retrieval mechanisms and generative models	Retrieval noise propagation, persistent hallucinations, high latency, absence of unified evaluation framework for grounding quality	Identifies retrieval noise as critical failure mode; motivates EvoStrategy's reconciliation gate that prevents noisy records from entering the analytics layer	2026
5	Gen AI Solution for Enhancing Hyper-Personalized Automotive Sales Proposal based on Customer Requirements	VAE-based latent representation model extracts 230+ attributes for personalized matching; Gemini 1.5 generates contextual proposals	Lack of integrated RAG framework increases hallucination risk; no cross-document verification mechanism for attribute consistency	Demonstrates generative AI's utility for structured business outputs; highlights need for verification gates when AI outputs feed decision pipelines	2025

S.no	Paper Title	Methodology	Limitations	Research Gaps Identified	Year
6	RAG vs GraphRAG: A Systematic Evaluation and Key Insights	Empirical comparison across multi-hop QA reasoning tasks with standardized performance metrics	Does not explore multimodal document ingestion or integration with numerical forecasting; restricted to text-only retrieval benchmarks	Empirical superiority of GraphRAG in multi-hop tasks validates knowledge graph approach for linking financial entities across document types	2025
7	Knowledge Graph Retrieval-Augmented Generation for LLM-based Recommendation	KG-RAG architecture combining knowledge graph retrieval with LLM generation for domain-specific recommendation tasks	KG utility demonstrated in recommendation domain; not yet applied to enterprise revenue simulation or financial document reconciliation	KG-RAG domain transfer potential directly applicable to IT financial knowledge graphs linking projects, vendors, and cost categories	2025
8	Mitigating Privacy Risks in RAG via Locally Private Entity Perturbation	Formal differential privacy and local privacy mechanisms applied to RAG retrieval pipelines	Provides privacy methods but does not address multi-hop graph reasoning or financial entity reconciliation across heterogeneous sources	Privacy-preserving RAG mechanisms validate the feasibility of sovereign local deployment for sensitive enterprise financial pipelines	2025
9	LLaVA-CoT: Let Vision Language Models Reason Step-by-Step	Vision-LLM chain-of-thought methods to improve stepwise visual reasoning and document understanding for complex layouts	Improves multimodal extraction but requires pipeline engineering for handwriting and enterprise document graphs	Step-by-step visual reasoning improves extraction from complex financial document layouts, reducing downstream reconciliation errors	2025
10	Construction of Intelligent Decision Support Systems through Integration of Retrieval-Augmented Generation and Knowledge Graphs	Framework IKEDS combining KGs and RAG for cross-domain decision support with structured knowledge retrieval	High manual overhead for knowledge graph maintenance; accuracy degrades in low-complexity decision tasks	IKEDS demonstrates RAG+KG decision support architecture; EvoStrategy extends this to verified financial forecasting by adding reconciliation gate	2025
11	Towards Robust Tampered Text Detection in Document Image: New Dataset and New Solution	DTD model with Frequency Perception Head and Multi-view Iterative Decoder on DocTamper dataset of 170,000 images; achieves F1=0.914	Requires high-resolution document scans; cross-domain performance degrades on non-Latin script documents	Document tamper detection can be integrated into EvoStrategy's ingestion stage to flag potentially fraudulent financial documents before reconciliation	2023
12	Enhancing Tampered Text Detection through Frequency Feature Fusion and Decomposition	FFDN combining Visual Enhancement Module and Wavelet-like Frequency Enhancement; achieves 8.58% F1 improvement over DTD on DocTamper	Wavelet decomposition adds computational overhead limiting real-time deployment on resource-constrained servers	Frequency-decomposition tamper detection provides deployable verification backbone for financial document authenticity checks in EvoStrategy	2024
13	RSTGCN: Railway-Centric Spatio-Temporal Graph Convolutional Network for Train Delay Prediction	Spatio-temporal GCN capturing delay propagation patterns across railway network nodes with temporal dependencies	Domain-specific to railway networks; temporal modeling techniques not directly transferable to irregular financial transaction timestamps	Temporal graph modeling techniques for propagation detection inform EvoStrategy's approach to tracing discrepancy propagation chains across linked documents	2025
14	Real-time Student Engagement Monitoring in Classroom Environments using Machine Learning and Computer Vision	ML and computer vision pipeline for real-time multi-person engagement detection with facial action unit analysis	Controlled environment evaluation; limited scalability to unstructured real-world monitoring without domain adaptation	Demonstrates production real-time ML pipeline architecture applicable to EvoStrategy's real-time document ingestion and discrepancy flagging workflow	2025
15	Generative AI Meets Large Language Models: Evolution, Challenges, and Future Directions	Comprehensive analysis of GenAI and LLM convergence trends evaluating architectural evolution and deployment challenges	Broad survey without domain-specific evaluation; limited coverage of financial document processing use cases	Identifies hallucination and data grounding as primary failure modes of generative AI; directly motivates EvoStrategy's reconciliation-gated forecasting approach	2025
16	Object Detection using Vision Transformer and Deep Learning for Computer Vision Applications	ViT-based object detection pipeline combining transformer attention with CNN feature extraction for multi-class detection	Computationally intensive; optimized for image classification rather than structured document layout understanding	Vision transformer attention mechanisms applicable to layout-aware document field extraction, improving structured data recovery from complex financial PDFs	2025
17	AI-driven Road Inspection with SUD-ROAD: High-resolution LiDAR Benchmark and Cross-dimensional Semantic Segmentation	Cross-dimensional semantic segmentation pipeline combining LiDAR point clouds with RGB imagery for road defect detection	Domain-specific; segmentation architecture not directly transferable to document region detection without significant retraining	Cross-modal fusion architecture demonstrates combining spatial and textual signals, analogous to EvoStrategy's fusion of OCR text with document layout features	2026
18	Crash Prediction under Limited CV Coverage: an Ensemble Deep Learning Model Integrating Multi-source Traffic Data	Ensemble deep learning combining CNN, LSTM, and GBM for crash prediction integrating heterogeneous traffic data sources	Requires dense sensor coverage; model interpretability limited for regulatory compliance requirements in safety-critical applications	Multi-source heterogeneous data fusion methodology directly applicable to EvoStrategy's integration of invoices, ledgers, and purchase orders from disparate formats	2026
19	ForgeFinder: Perceptive Multimodal Deepfake Detection via Multi-grained Forgery Localization	Multimodal deepfake detection with multi-grained forgery localization across visual and textual modalities	High computational cost; primarily evaluated on face deepfakes rather than document-level text manipulation	Multi-grained localization approach applicable to identifying the specific field within a financial document that has been altered, improving reconciliation	2026
20	Advancing Cybersecurity: AI-Driven Computer Vision and Machine Learning Models for Real-Time Threat Detection	Real-time AI pipeline combining computer vision and ML for multi-vector cybersecurity threat detection and prevention	Computational overhead for real-time inference; threat taxonomy may not generalize across organizational security perimeters	Automated anomaly detection architecture applicable to EvoStrategy's discrepancy detection engine, adapting threat-detection pattern recognition to financial inconsistency identification	2026

III. PRIOR ART SEARCH

1) Document OCR and Structured Extraction

What Exists: Tesseract 5.x (LSTM-based) and PaddleOCR 2.x achieve >95% character accuracy on clean printed text. Layout-aware models (LayoutLM, PaddleOCR's structure recognition) handle multi-column tables and form fields.

Gap: Accuracy degrades significantly on low-quality scans, handwritten annotations, and non-standard financial document layouts. Post-extraction validation — verifying that extracted amounts match document-level checksums or totals — is rarely implemented in published systems.

2) Entity Resolution and Record Linkage

What Exists: String normalization (lowercasing, punctuation removal), fuzzy matching (Levenshtein, Jaro-Winkler), and blocking-based record linkage are well-established techniques for deduplication and entity resolution in databases.

Gap: Published entity resolution systems operate on structured database records. Applying entity resolution to heterogeneous financial documents — where the same vendor may appear in different scripts, abbreviations, or with typographical errors — remains an under-evaluated problem with no standard benchmark.

3) Cross-Document Financial Reconciliation

What Exists: Accounting software implements within-system reconciliation (e.g., bank statement vs. ledger). Some research on table-to-table numerical consistency checking exists in the scientific data domain.

Gap: No widely published system performs automated cross-document reconciliation across heterogeneous unstructured documents (PDF invoices, scanned receipts, Excel ledgers, email embedded payment records) with automatic escalation for human review. This is the primary contribution of EvoStrategy.

4) Verified Financial Forecasting

What Exists: Strong time-series forecasting libraries (Prophet, statsmodels, scikit-learn) provide robust regression and ARIMA implementations. Business forecasting tools integrate with structured data warehouses.

Gap: Existing forecasting systems do not distinguish between verified and unverified input data. They apply the same models regardless of data quality, producing projections that inherit the errors in the underlying records. EvoStrategy's reconciliation-gated forecasting architecture addresses this gap.

5) Sovereign and Offline AI Deployment

What Exists: Offline RAG deployments using quantized local models (Ollama, LM Studio) and local vector

databases (Chroma, FAISS) demonstrate that AI document pipelines can operate without cloud connectivity.

Gap: Comprehensive, evaluated deployments combining document ingestion, cross-document reconciliation, human-in-the-loop verification, and forecasting within a single locally-deployed pipeline are not available in published literature. EvoStrategy targets this gap specifically for the IT services financial management domain.

IV. GAPS IDENTIFIED THROUGH LITERATURE SURVEY


- No existing system treats reconciliation as a prerequisite gate for analytics. Current tools visualize and forecast from whatever data is available, including erroneous data.
- Entity resolution for financial documents remains an unsolved benchmark problem. Vendor alias matching across heterogeneous document formats lacks standardized evaluation datasets.
- Human-in-the-loop escalation is poorly addressed. Most automated reconciliation systems either fully automate (and produce silent errors) or fully rely on humans (and are slow). The hybrid escalation architecture is underexplored.
- Cross-document numerical linking — specifically, detecting when figures in document A should match figures in document B — has no established methodology for unstructured financial documents.
- Sovereign deployment of complete financial intelligence pipelines (ingestion → reconciliation → verification → forecasting) without cloud dependency has not been systematically demonstrated.

V. PROBLEM FORMULATION

The problem addressed in this work is to develop an automated and auditable mechanism for reconciling financial information distributed across heterogeneous enterprise documents. Given a set of financial documents such as invoices, purchase orders, payment memos, and expense ledgers, the system must extract relevant transaction attributes, establish relationship between corresponding records, and identify inconsistencies such as numerical mismatches, entity aliases, duplicate entries, and missing records. The system should automatically resolve discrepancies within predefined tolerance limits while escalating unresolved cases for human verification. Furthermore, only reconciled and verified records should be permitted to enter the downstream analytics and forecasting processes, ensuring that financial projections are traceable to validated source information. The complete process must operate locally to preserve the confidentiality and sovereignty of enterprise financial data.

VI. METHODOLOGY

EvoStrategy is implemented as a four-stage, locally deployed document intelligence pipeline consisting of Document Ingestion and Extraction, Cross-Document Reconciliation, Human-in-the-Loop Verification, and Verified Analytics and Forecasting. The overall architecture follows a reconciliation-gated approach in which only verified records are permitted to enter downstream analytics and forecasting.




A) Document Ingestion and Extraction

The system accepts heterogeneous financial documents including PDFs, scanned images, and spreadsheet records. PDF documents are rendered using PyMuPDF, while OpenCV is used for preprocessing operations such as deskewing, denoising, and contrast normalization. Tesseract OCR is used for printed text extraction and PaddleOCR for tables and complex layouts. The extracted information is converted into structured transaction records containing fields such as vendor, date, amount, category, document type, source file, and confidence score. Low-confidence extractions are flagged for subsequent processing.

B) Cross-Document Reconciliation

The extracted records are linked according to predefined relationships such as invoice–payment memo, invoice–purchase order, and invoice–ledger entries. Vendor entities are resolved using string normalization, spaCy-based entity extraction, and RapidFuzz similarity matching. Related documents are linked using vendor identity and date proximity, after which numerical values are compared using predefined tolerance limits. Discrepancies within the tolerance range are automatically resolved, while larger or ambiguous discrepancies are classified and escalated for review.



C) Human-in-the-Loop Verification

Discrepancies that cannot be reliably solved automatically are presented through a Streamlit-based verification dashboard. Reviewers can inspect the corresponding source documents and extracted values side-by-side and perform ACCEPT, REJECT, or CORRECT actions. Each decision is recorded with reviewer identity and timestamp, providing an auditable history. Verified records are promoted to the reconciled record store, while rejected or corrected records are retained with their corresponding decision information.

D) Verified Analytics and Forecasting

Only matched, automatically resolved, and human-verified records are passed to the analytics layer. The system generates spending trends, category-level budget consumption, and runway projections using linear regression and ARIMA-based forecasting. A What-If simulation module allows users to modify selected financial parameters and observe their impact on projected outcomes. The resulting visualizations include source traceability information such as record count, date range, and verification status.

The complete pipeline is deployed locally using Python, SQLite, Streamlit, and Docker-based components, ensuring that enterprise financial documents remain within the organization's infrastructure. This architecture establishes the reconciled record store as the trusted data source for downstream financial intelligence.

IX. CONCLUSION

This paper presented EvoStrategy, a unified enterprise intelligence framework that establishes document reconciliation as a prerequisite for reliable financial analytics rather than an afterthought. By combining OCR-based ingestion, cross-document entity and numerical reconciliation, and human-in-the-loop verification within a single locally deployed pipeline, the proposed system addresses a gap left open by existing extraction and forecasting tools: the failure to distinguish verified data from unverified data before it reaches downstream analysis. The reconciliation-gated architecture, together with complete audit traceability, provides a foundation for generating revenue forecasts and What-If simulations that enterprises can trust and trace back to their original source documents, while preserving data confidentiality through fully local deployment.